



# Introduction to Joins



a join shows a result set, containing fields derived from two or more tables

Error code:1175 workaround:

74 14:03:21 DELETE FROM dept\_manager\_dup WHERE dept\_no = 'd001'

Error Code: 1175. You are using safe update mode and you tried to update a table without a WHERE that uses a KEY column. To disable safe mode, toggle the option in Preferences -> SQL Editor and reconnect.

Why the error happens:

MySQL Workbench's *safe update mode* blocks any DELETE or UPDATE query unless:

- The WHERE clause uses a key column (usually a PRIMARY KEY or an indexed column), or
- You include a LIMIT.

In our table, dept\_no is NOT a key column (because our CREATE TABLE did not define a PRIMARY KEY).

Therefore:

```
DELETE FROM dept_manager_dup
WHERE dept_no = 'd001';
```

gets blocked.

This will delete all rows matching dept\_no = 'd001' while staying within safe-update rules.

#### Alternative workaround (without disabling safe mode):

Use the primary key — emp\_no — in the WHERE clause:

```
DELETE FROM dept_manager_dup
WHERE emp_no IN (
    SELECT emp_no FROM dept_manager WHERE dept_no = 'd001'
);
```

However, MySQL won't allow deleting from a table while selecting from the same table unless wrapped differently.

So the safer practical workaround is:

✓ Use LIMIT.

DON'T DISABLE SAFE MODE.

Final recommendation:

```
DELETE FROM dept_manager_dup
WHERE dept_no = 'd001'
LIMIT 1000;
```

## INNER JOIN



```
SQL   SELECT
        table_1.column_name(s), table_2.column_name(s)
FROM
        table_1
JOIN
        table_2 ON table_1.column_name = table_2.column_name;
```

Lecturer notes: In a select statement, write all columns we wish to see in the result. It is essential to designate the tables to which the columns belong, as the data is not contained in a single table this time. That's why, besides typing the keyword "from", and the name of the first table, we should proceed by writing "join" and the name of the second table. The syntax allows us to specify the fields we would like to see in the result in the tables we are matching.

## ● INNER JOIN



```
SELECT  
    t1.column_name, t1.column_name, ..., t2.column_name, ...  
FROM  
    table_1 t1  
JOIN  
    table_2 t2 ON t1.column_name = t2.column_name;
```

## aliases

Lecturer notes: A fundamental coding practice that professionals use in the joint syntax. [Aliases](#).

More precisely, we're talking about aliases of the table's names. This means table one can be renamed to say T1 and table two to T2. When used for assigning table names, the aliases are usually added right after the original table name, without using the keyword "as". Then, instead of typing the entire table's names in the select block, we can use T1 and T2, respectively.

The screenshot shows a SQL management interface with a tree view on the left and a results grid on the right. The results grid displays three rows of data from an INNER JOIN query:

1	006	110000	Quality Management
2	006	110765	Quality Management
3	006	110854	Quality Management

Two blue callout boxes highlight key points about INNER JOINs:


- A top box states: "inner joins extract only records in which the values in the related columns match"
- A bottom box states: "null values, or values appearing in just one of the two tables and not appearing in the other, are not displayed"

The status bar at the bottom indicates the query was run "Read Only".

So when using the JOIN/ INNER JOIN function, we won't get any null values or incomplete records.

## INNER JOIN

- [And what if such matching values did not exist?](#)



Simply, the result set will be empty. There will be no link between the two tables.

Lecturer notes: The terms "JOIN" and "INNER JOIN" refer to the same type of SQL operation used to combine rows from two or more tables based on a related column between them. Here are the key points about their usage:

1. **Functionality:** Both "JOIN" and "INNER JOIN" return rows where there is a match in both tables. If there are no matching records, those rows will not be included in the result set.
2. **Interchangeability:** As per the course material, they are functionally equivalent, meaning you can use either term without affecting the outcome of your SQL query.
3. **Readability:** Using "INNER JOIN" may enhance clarity, especially in queries that involve multiple types of joins (like LEFT JOIN or RIGHT JOIN). This helps in identifying which type of join is being applied at a glance.
4. **Preference:** While you can use either "JOIN" or "INNER JOIN," some developers prefer to use "INNER JOIN" for the sake of clarity, particularly in complex queries.

My work:


```

200 #BIG BOY QUESTION: Extract a list containing information about all managers' employee number,
201 -- first and last name, department number, and hire date.
202 • SELECT
203     e.emp_no, e.first_name, e.last_name, dm.dept_no, e.hire_date
204 FROM
205     employees e #`employees` is given the alias 'e' to make references shorter and clearer so in the SELECT 'e.emp_no' comes from the employees table ar
206     INNER JOIN
207     dept_manager dm #'dept_manager' is given the alias 'dm' for the same reason - THIS METHOD ELIMINATES THE USE OF 'AS' function. we do this when joini
208     ON e.emp_no = dm.emp_no; #The ON clause specifies how the two tables relate:
209             -- both tables contain a column called emp_no, which is the primary link between them,
210             -- This means each manager must exist in the employees table,
211             -- and the JOIN returns only the rows where emp_no appears in BOTH tables.

```

Result Grid					
	emp_no	first_name	last_name	dept_no	hire_date
▶	110022	Margareta	Markovitch	d001	1985-01-01
	110039	Vishwani	Minakawa	d001	1986-04-12
	110085	Ebru	Alpin	d002	1985-01-01
	110114	Isamu	Legleitner	d002	1985-01-14
	110183	Shirish	Ossenbruggen	d003	1985-01-01
	110228	Karsten	Sigstam	d003	1985-08-04
	110303	Krassimir	Wegerle	d004	1985-01-01
	110344	Rosine	Cools	d004	1985-11-22
	110386	Shem	Kieras	d004	1988-10-14
	110420	Oscar	Ghazalie	d004	1992-02-05
	110511	DeForest	Hagimont	d005	1985-01-01

## LEFT JOIN



The Venn diagram we see here allows us to visualise how a Left Join works. Its output allows us to see all records from the table on the left side of the Join, including all matching rows of the two tables. That's why, compared to the Inner Join, the results set, coloured in red, includes the rest of the area of the left table. In SQL terms, this translates to retrieving all matching values of the two tables, plus all values from the left table that match no values from the right table.

```

SELECT
    t1.column_name, t1.column_name, ..., t2.column_name, ...
FROM
    table_1 t1
LEFT JOIN
    table_2 t2 ON t1.column_name = t2.column_name;

```

```

14
15
16 # LEFT JOIN
17 • SELECT
18     m.dept_no, m.emp_no, d.dept_name
19     FROM
20         dept_manager_dup m
21             LEFT JOIN
22                 departments_dup d ON m.dept_no = d.dept_no
23     GROUP BY m.emp_no
24     ORDER BY m.dept_no;
25

```

26 rows (left join)			20 rows (inner join)		
dept_no	emp_no	dept_name	dept_no	emp_no	dept_name
d001	999905	NULL	d003	110228	Human Resources
d001	999907	NULL	d003	110183	Human Resources
d001	999904	NULL	d004	110344	Production
d001	999905	NULL	d004	110420	Production
d002	110085	NULL	d004	110303	Production
d002	110114	NULL	d004	110386	Production
d003	110183	Human Resources	d005	110567	Development
d003	110228	Human Resources	d005	110511	Development
d004	110420	Production	d006	110800	Quality Management
d004	110303	Production	d006	110765	Quality Management
d004	110386	Production	d006	110854	Quality Management
d004	110344	Production	d006	110725	Quality Management
d005	110511	Development	d007	111035	Sales
d005	110567	Development	d007	111133	Sales
d006	110765	Quality Management	d008	111400	Research
d006	110854	Quality Management	d008	111534	Research
d006	110725	Quality Management	d009	111784	Customer Service
d006	110800	Quality Management	d009	111939	Customer Service
d007	111133	Sales	d009	111692	Customer Service
d007	111035	Sales	d009	111877	Customer Service
d008	111534	Research			
d008	111400	Research			
d009	111692	Customer Service			
d009	111877	Customer Service			
d009	111784	Customer Service			
d009	111939	Customer Service			

Lecturer notes: It returned 26 rows. Six rows more than the 20 rows we obtained in the example about Inner Joins. Basically, this is proof that, unlike what we saw for Inner Joins, when working with Left Joins, the order in which you join tables matters. Having the manager's table, M, or the department's table, D, on the left can change results completely.

e.g.

```

1 • SELECT
2   d.dept_no, m.emp_no, d.dept_name
3   FROM
4     departments_dup d
5   LEFT JOIN
6     dept_manager_dup m ON m.dept_no = d.dept_no
7   ORDER BY d.dept_no;
8

```

**departments\_dup d**

dept_no	emp_no	dept_name
NULL	NULL	Public Relations
d001	NULL	Marketing
d003	110383	Human Resources
d004	110228	Human Resources
d004	110344	Production
d004	110420	Production
d004	110303	Production
d004	110386	Production
d005	110567	Development
d005	110511	Development
d006	110690	Quality Management
d007	110728	Sales
d008	110745	Research
d009	110759	Customer Service
d010	110761	IT
d011	110762	IT

**dept\_manager\_dup m**

dept_no	dept_name
d001	Public Relations
d003	Marketing
d003	Human Resources
d004	Production
d004	Production
d004	Production
d005	Development
d006	Quality Management
d007	Sales
d008	Research
d009	Customer Service
d010	IT
d011	IT

Action Output

- Time Action
  - 2 15:37:24 SELECT m.dept\_no, m.emp\_no, d.dept\_name FROM departments\_dup d ... 26 rows(s) returned
  - 3 15:38:30 SELECT d.dept\_no, m.emp\_no, d.dept\_name FROM departments\_dup d ... 26 rows(s) returned

Duration / Fetch  
0.000 sec / 0.000 sec

365 Careers

LEFT JOIN = LEFT OUTER JOIN (Used interchangeably)

## RIGHT JOIN

### RIGHT JOIN

their functionality is identical to LEFT JOINs, with the only difference being that the direction of the operation is inverted

• LEFT and RIGHT joins are perfect examples of one-to-many relationships

1 manager

1 department

dept\_manager\_dup

dept\_no CHAR(4)  
emp\_no INT  
from\_date DATE  
to\_date DATE

departments\_dup

dept\_no CHAR(4)  
dept\_name VARCHAR(40)

Lecturer notes: In addition, when talking about relationships, left and right joins are perfect examples of one-to-many relationships in my SQL. For instance, in our last example, when we used a left join, each department from the department's duplicate table, as represented by the department number, could have been the department of one or more managers from the department manager duplicate table. A manager who is also an employee can belong to a single department only. This is an example of how the one-to-many relationship can be exhibited in a left or right join case.


The same results we obtained by using the JOIN function can also be obtained via the WHERE function: **THE NEW & OLD JOIN SYNTAX**

### The New and the Old Join Syntax

- WHERE (the Old Join Syntax)

 SQL

```
SELECT
    t1.column_name, t1.column_name, ..., t2.column_name, ...
FROM
    table_1 t1,
    table_2 t2
WHERE
    t1.column_name = t2.column_name;
```



The screenshot shows a MySQL Workbench interface. The query editor window contains the following code:

```
11 -- WHERE
12 • SELECT
13     m.dept_no, m.emp_no, d.dept_name
14 FROM
15     dept_manager_dup m,
16     departments_dup d
17 WHERE
18     m.dept_no = d.dept_no
19 ORDER BY m.dept_no;
```

The results grid displays the following data:

dept_no	emp_no	dept_name
d003	110228	Human Resources
d003	110183	Human Resources
d004	110344	Production
d004	110420	Production
d004	110303	Production
d004	110386	Production
d005	110511	Development
d005	110567	Development
d006	110725	Quality Management

## The New and the Old Join Syntax

- **JOIN or WHERE?**

- the retrieved output is *identical*
- using WHERE is more *time-consuming*
- the WHERE syntax is perceived as *morally old* and is rarely employed by professionals
- the JOIN syntax allows you to modify the connection between tables easily

## CROSS JOIN FUNCTION

### CROSS JOIN

- a cross join will take the values from a certain table and connect them with all the values from the tables we want to join it with
- **INNER JOIN**
  - typically connects *only the matching values*
- **CROSS JOIN**
  - connects *all the values, not just those that match*
  - the *Cartesian product of the values of two or more sets*

Lecturer notes: A cross join will take the values from a certain table and connect them with all the values from the tables we want to join it with. This is in contrast to the inner join that typically connects only to matching values. A cross join will connect all the values, not just those that match. That's why, from a mathematical point of view, a cross join is the Cartesian product of the values of two or more sets.

A cross join is particularly useful when the tables in a database are not well-connected. We must admit that the Employees Database is not really suitable for applying this kind of join meaningfully since its tables are indeed well connected. However, we can still use the employees database just to do an exercise with a cross join, can't we?

Here's an example:

The screenshot shows a SQL query window titled "CROSS JOIN". The query itself is:

```
1 • SELECT
2     dm.* , d.*
3     FROM
4         dept_manager dm
5     CROSS JOIN
6         departments d
7     ORDER BY dm.emp_no , d.dept_no; ]
```

To visualise our output better, we will order the values by employee number as specified in the department manager table. and then by department number as specified in the department's table.

RESULTS:

The screenshot shows a MySQL Workbench interface with a result grid. The grid has columns: emp\_no, dept\_no, from\_date, to\_date, dept\_no, and dept\_name. The data consists of 14 rows, where each row from the first table (dept\_manager) is paired with every row from the second table (departments). The dept\_name column lists Finance, Human Resources, Production, Development, Quality Management, Sales, Research, Customer Service, Marketing, and Finance again.

emp_no	dept_no	from_date	to_date	dept_no	dept_name
110039	d001	1991-10-01	9999-01-01	d002	Finance
110039	d001	1991-10-01	9999-01-01	d003	Human Resources
110039	d001	1991-10-01	9999-01-01	d004	Production
110039	d001	1991-10-01	9999-01-01	d005	Development
110039	d001	1991-10-01	9999-01-01	d006	Quality Management
110039	d001	1991-10-01	9999-01-01	d007	Sales
110039	d001	1991-10-01	9999-01-01	d008	Research
110039	d001	1991-10-01	9999-01-01	d009	Customer Service
110085	d002	1985-01-01	1989-12-17	d001	Marketing
110085	d002	1985-01-01	1989-12-17	d002	Finance

We can observe that all department managers have been connected with all departments. In other words, nine different department numbers correspond to the employee number of each manager.

NOTICE: how emp\_no and dept\_no are in consecutive order? Remember it was initially ordered by emp\_no, then it would be ordered by dept\_no for the second table. And in the joining of both tables, where the first table ends, a new order is established.

### ANOTHER INTERESTING WAY OF DOING A CROSS JOIN -WITHOUT WHERE OR THE JOIN STATEMENT:

The screenshot shows a MySQL Workbench interface with a query editor and a result grid. The query uses a Common Table Expression (CTE) named 'CROSS JOIN' to perform a cross join between 'dept\_manager' and 'departments' tables. The result grid shows the same 14 rows as the previous screenshot, where each row from 'dept\_manager' is paired with every row from 'departments'. A large blue bracket highlights the CTE part of the query, and another blue bracket highlights the result grid.

emp_no	dept_no	from_date	to_date	dept_no	dept_name
110085	d002	1985-01-01	1989-12-17	d001	Marketing
110085	d002	1985-01-01	1989-12-17	d002	Finance
110085	d002	1985-01-01	1989-12-17	d003	Human Resources
110085	d002	1985-01-01	1989-12-17	d004	Production
110085	d002	1985-01-01	1989-12-17	d005	Development
110085	d002	1985-01-01	1989-12-17	d006	Quality Management
110085	d002	1985-01-01	1989-12-17	d007	Sales
110085	d002	1985-01-01	1989-12-17	d008	Research
110085	d002	1985-01-01	1989-12-17	d009	Customer Service
110085	d002	1989-12-17	9999-01-01	d001	Marketing

Well, the result is the same. The answer is that this is exactly the output of a join between these two tables with no wear statement with which we can set a condition to the tables. Hence, the result is a cross join between department manager and departments.

## GOING Further in our analysis:

The screenshot shows a MySQL Workbench interface with a query editor and two result grids. The query editor contains the following SQL code:

```

CROSS JOIN x
14 ORDER BY dm.emp_no , d.dept_no;
15
16 • SELECT
17     dm.* , d.*
18     FROM
19         dept_manager dm
20             JOIN
21                 departments d
22     ORDER BY dm.emp_no , d.dept_no;
  
```

The result grid on the left shows the output of the query, which is a Cartesian product of the two tables. The result grid on the right shows the individual tables: 'dept\_manager' and 'departments'. A large blue 'X' is drawn over the 'departments' table, indicating it is not part of the current join operation. Below the tables, the text 'CROSS JOIN' is written in bold.

Lecturer notes: We can rewrite the previous example in a third way, like this. See, there is no sign of the word cross in this query. Although the result is the same as before. In addition, we don't have a conditional statement connecting the two tables, neither in a WHERE statement nor after the ON keyword. Nevertheless, MySQL will interpret this join as a cross join and won't raise a syntax error. You can even write it as an inner join, and the result will still be the same because no condition has been assigned. The truth is that writing an inner join without the keyword ON is not considered best practice. Writing a cross join, on the other hand, will help your colleagues have a much clearer idea about the expected result while reading your code. That's why my SQL is so powerful. Often, there are many ways that could lead you to an identical result. But of course, clarity is a substantial part of writing good code. Hence, in this course, we stick to using best practices only.

The screenshot shows a MySQL Workbench interface with a message box overlaid on the query editor. The message box contains two statements:

- JOIN without ON = not considered best practice**
- CROSS JOIN = best practice**

The message box is blue with white text. The background shows the MySQL Workbench interface with a query editor containing:

```

17 FROM
18
19 dept_manager dm
  
```

The status bar at the bottom shows the query: "2 11:14:17 SELECT dm.\* , d.\* FROM dept\_manager dm, departments d ORDER BY dm.emp\_no , d.dept\_no; 216 rows(s) returned".

However, what should we do if we want to display all departments, but the one where the manager is currently the head of? To be frank, there's nothing simpler than that. All we have to do is add a where clause containing the condition that the department number and the department's table is different from the department number of the employee in question:

CROSS JOIN\*

```

24 • SELECT dm.* , d.*
25   dm.* , d.*
26   FROM departments d
27     CROSS JOIN
28   dept_manager dm
29 WHERE d.dept_no <> dm.dept_no
30 ORDER BY dm.emp_no , d.dept_no;

```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

emp_no	dept_no	from_date	to_date	dept_no	dept_name
110035	d002	1985-01-01	1989-12-17	d001	Marketing
110035	d002	1985-01-01	1989-12-17	d003	Human Resources
110035	d002	1985-01-01	1989-12-17	d004	Production
110035	d002	1985-01-01	1989-12-17	d005	Development
110035	d002	1985-01-01	1989-12-17	d006	Quality Management
110035	d002	1985-01-01	1989-12-17	d007	Sales
110035	d002	1985-01-01	1989-12-17	d008	Research
110035	d002	1985-01-01	1989-12-17	d009	Customer Service
110039	d001	1989-12-17	9999-01-01	d002	Marketing
110039	d002	1989-12-17	9999-01-01	d003	Human Resources

Result 5 × Read Only

Output:

Action Output

#	Time	Action	Message	Duration / Fetch
3	11:14:55	SELECT dm.* , d.* FROM dept_manager dm	JOIN departments d ... 216 row(s) returned	0.000 sec / 0.000 sec
4	11:15:20	SELECT dm.* , d.* FROM dept_manager dm	INNER JOIN departments d ... 216 row(s) returned	0.000 sec / 0.000 sec
5	11:16:33	SELECT dm.* , d.* FROM departments d	CROSS JOIN dept_ma... 192 row(s) returned	0.000 sec / 0.000 sec

365 Careers Academy

Lecturer notes: After executing this query, we know it is right because we can see that the department in which the manager is working has not been shown. Moreover, if we compare the number of records retrieved here with the ones retrieved in the last example, the difference will be exactly the number of managers in the department manager's table.

Finally, we can cross-join more than two tables. However, we should be really careful when doing so because if the tables contain a lot of records, there is a chance the result might become too big. And hence MySQL won't be able to execute the query. This problem may arise if you are cross-joining two tables containing lots of records as well. Nevertheless, when the tables do not contain too many records, cross joins can become the perfect tool you need.

Let's make a cross join and combine it with the good old inner join:

CROSS JOIN\*

```

23
24 • SELECT e.* , d.*
25   e.* , d.*
26   FROM departments d
27     CROSS JOIN
28   dept_manager dm
29   JOIN
30   employees e ON dm.emp_no = e.emp_no

```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

emp_no	birth_date	first_name	last_name	gender	hire_date	dept_no	dept_name
110022	1956-09-12	Maroareta	Markovitch	M	1985-01-01	d002	Finance
110022	1956-09-12	Maroareta	Markovitch	M	1985-01-01	d003	Human Resources
110022	1956-09-12	Maroareta	Markovitch	M	1985-01-01	d004	Production
110022	1956-09-12	Maroareta	Markovitch	M	1985-01-01	d005	Development
110022	1956-09-12	Maroareta	Markovitch	M	1985-01-01	d006	Quality Management
110022	1956-09-12	Maroareta	Markovitch	M	1985-01-01	d007	Sales
110022	1956-09-12	Maroareta	Markovitch	M	1985-01-01	d008	Research
110022	1956-09-12	Maroareta	Markovitch	M	1985-01-01	d009	Customer Service
110039	1963-06-21	Vishwani	Minakawa	M	1986-04-12	d002	Finance
110039	1963-06-21	Vishwani	Minakawa	M	1986-04-12	d003	Human Resources

Result 6 × Read Only

Output:

Action Output

#	Time	Action	Message	Duration / Fetch
4	11:15:20	SELECT dm.* , d.* FROM dept_manager dm	INNER JOIN departments d ... 216 row(s) returned	0.000 sec / 0.000 sec
5	11:16:33	SELECT dm.* , d.* FROM departments d	CROSS JOIN dept_ma... 192 row(s) returned	0.015 sec / 0.000 sec
6	11:18:48	SELECT e.* , d.* FROM departments d	CROSS JOIN dept_ma... 192 row(s) returned	0.015 sec / 0.000 sec

365 Careers Academy